

System & Software Security

Software Bills of Materials

Hands-On Project

| | |
|---------------|----------|
| Davide Baggio | s4426428 |
| João Pereira | s4443322 |
| Nuno Pereira | s4443330 |

November 2024

We confirm that this report was fully produced by the team members **Davide Baggio, João Pereira and Nuno Pereira** and we are jointly responsible for all content presented in this work. All used sources were attributed properly.

Contents

| | | |
|----------|-------------------------------|----------|
| 1 | Introduction | 1 |
| 2 | Dataset | 1 |
| 2.1 | Standards | 1 |
| 2.1.1 | CycloneDX | 1 |
| 2.1.2 | SPDX | 2 |
| 2.1.3 | SWID | 2 |
| 2.2 | Tools | 2 |
| 2.2.1 | Other tools | 3 |
| 2.3 | Repositories | 3 |
| 3 | Methodology | 4 |
| 4 | Results and Discussion | 4 |
| 4.1 | CycloneDX | 4 |
| 4.2 | SPDX | 5 |
| 4.3 | SWID | 6 |
| 5 | Conclusion | 6 |
| | References | 7 |
| A | Appendix A | 9 |
| B | Appendix B | 9 |

1 Introduction

Software Chain Security is an increasingly important challenge to tackle as early as possible in the Software Development process due to the catastrophic effect that vulnerabilities or security issues in downstream dependencies can cause to any software product, like in the SolarWinds [6] and Log4J [5] cases. It is therefore crucial from a development standpoint to be aware of what components are included in any software project and the potential vulnerabilities that they might introduce. Many solutions have been developed to aid in that process, the prime example being *package-managers* [1], tools that help managing software dependencies and version conflicts between dependencies. Examples of such package managers are NPM (Node Package Manager) and Cargo [24, 26].

Another solution to dependency and vulnerability tracking are *Software Bills of Materials* (SBOMs) [3], detailed listings of dependencies, their relations, licensing information and other metadata pertinent to software products. These provide a standard format to work and process dependency information, allowing easier communication and shareability. Currently, 3 mainstream SBOM standards exist: OWASP’s *CycloneDX*, Linux Foundation’s *SPDX* and NIST’s *SWID* [12, 13, 22].

In previous work [29], the authors performed a critical comparison on 5 articles from the literature [7–11] based on author-defined metrics. The authors provide insights on the state-of-the-art regarding SBOMs and provide guidelines for SBOM adoption and development, as well as future work that can be done in further research.

In an attempt to complement [29], this article reports on a hands-on comparison of the three mainstream SBOM standards by making use of available tools for each standard. SBOM generation tools for each standard were used on a set of 3 major Open-Source repositories found on GitHub and the resulting SBOM output files were compared, both between standards as well as between the tools of each standard.

2 Dataset

In this section, we detail the dataset used to perform the hands-on comparison between SBOM standards. In 2.1 we discuss the different available standards and the tools chosen to generate SBOMs for each one of them are described in 2.2. Based on the tools picked we chose, in 2.3, the repositories for which the SBOMs will be generated.

2.1 Standards

Each one of the three standard formats focuses on a specific part of the software supply chain, which can be reflected in the (meta)data that each standard stores and processes. The tools developed for each each standard also reflect these decisions.

2.1.1 CycloneDX

CycloneDX [13] is an SBOM standard format developed by the CycloneDX Core Working Group and backed by the OWASP Foundation with a focus on “cyber-risk reduction” [13] and security [7]. The standard supports writing BOMs for several domains of software development, such as Software BOMs (SBOMs), Cryptographic BOMs (CBOMs), Software-as-a-Service BOMs (SaaS BOMs), among others. Over 200 tools related to CycloneDX’s SBOM format are available at [CycloneDX’s official tool webpage](#).

For this hands-on comparison, we limited our search to *Open-Source* tools as these are free to access and use. Out of 172 listed Open-Source tools, 3 were chosen: CycloneDX `cdxgen` [14], `build-info-go` [21] and `syft` [17]. `syft` also supports SPDX, which will be discussed in 2.1.2.

Other tools exist but they are either unrelated (SBOM analysis, VEX generation, ...), too specific (official SBOM generators for several existing programming languages and build tools) or too limited on the supported development environments.

2.1.2 SPDX

SPDX (System Package Data Exchange) [22] is a standard format maintained by the Linux Foundation for communicating software bill of material information, including provenance, license, security, AI and other related information. All these characteristics make it a very versatile tool for software supply chain management. Nowadays it is used in the Linux kernel and in many package managers.

Official SPDX tools include an online SBOM inspector, validator and converter [28], build tools for Java (Gradle and Maven plugins) [19, 20], and several libraries for SBOM manipulation in different languages. Additionally, there are community-lead projects such as `spdx-sbom-generator`, and `syft` [17, 18] that aim to provide automatic SBOM generation for several different languages and package managers.

2.1.3 SWID

Software Identification (SWID) Tags [12] are a standard format for identifying software components and metadata, which can be used to generate SBOMs. Nowadays, the current standard of it is ISO/IEC 19770-2:2015 [2] and is maintained by the ISO/IEC JTC 1/SC 7/WG 21 committee [16].

A generated SWID Tag document consists of a well-organized collection of data fields that specify the software product, its version, identify the organizations and people involved in its creation and distribution, list the components that make up the software, define relationships between different software products and include additional metadata for further description. SWID tagging differs from CycloneDX and SPDX in that it is not a full-fledged SBOM format, as it doesn't aggregate information of all softwares so it's rather a standard for identifying software components and their metadata.

2.2 Tools

The following tools were used to generate SBOMs for the different standards:

cdxgen is an official tool developed and released on GitHub by the CycloneDX team, built around the idea of being a "polyglot SBOM generator that is user friendly, precise and comprehensive". It provides a comprehensive SBOM generator for different versions of the CycloneDX Standard. It also provides a Server mode, automatic licensing information, and Docker/OCI container support. It is distributed as an NPM package.

build-info-go is a CLI tool to generate BuildInfo metadata, a custom format designed to encapsulate software components, their versions and their dependencies. The tool supports multiple languages and package managers and has the option to export the produced BuildInfo output into a valid CycloneDX JSON file.

syft is a "CLI tool and Go library for generating a Software Bill of Materials (SBOM) from container images and filesystems" [17]. It provides visibility into vulnerabilities and license compliance and can interact with modern vulnerability scanners such as Grype. It can output information in over 10 output formats, including custom user-defined formats specified by templates, and supports over 20 ecosystems.

spdx-sbom-generator is a tool developed by the SPDX community to generate SPDX SBOMs for several different languages and package managers. This tool is still in its infancy, while it boasts support for a lot of languages, it is not battle tested and crashed with most of what we tried. It also does not have good enough configuration.

2.2.1 Other tools

These tools were tried but we were not able to generate SBOMs using them:

spdx-gradle-plugin is the official Gradle plugin for SPDX SBOM generation, which needs to be configured in the project's `build.gradle` file. This requires intimate knowledge of the repositories build process, which we do not have. Due to this and an overall lack of documentation and user friendliness, this tool was not used.

swid-builder a Java API for building SWID and CoSWID tags. This library provides a set of builder patterns that can be used together to generate tags, however it is not a standalone tool and the provided [documentation](#) is not very clear on how to use it. Thus, it is not suitable for our use case.

swid-maven-plugin is a Maven plugin published by NIST [23] that generates SWID Tags for Java projects and is compliant with the above mentioned ISO standard. It needs to be configured in the project's `pom.xml` file and additionally it requires the assembly descriptor in `src/assembly/bin.xml` to be configured as well. The plugin is not maintained anymore but still produces a valid SWID Tag file.

swid-generator as alternatives to the official tools, we also looked for unofficial ones on GitHub and found a tool developed by Labs64 [4], which also supports Gradle projects. The downside of it is that it requires implementing a custom portion of code in the project to generate the SWID Tags, which is not as straightforward as the other tool and it might be unsuitable for complex projects.

By the same authors, there is also a Maven plugin which is unfortunately not maintained anymore and uses an outdated version of Java, which results in a lot of errors when trying to build it, as testified by [this issue](#).

2.3 Repositories

To ensure a fair comparison between standards, we chose a representative set of major Open-Source repositories that could be analyzed by most, if not all, of the tools selected.

As such, we have picked 3 repositories from GitHub:

- Apache Kafka [15] is a modern event-streaming platform written in Java by the Apache Software Foundation. Initially developed at LinkedIn to accommodate their

growing message processing needs, it now serves as the backbone for many asynchronous, event-driven and streaming systems around the world.

- `numpy` [25] is a Python library that provides a "multi-dimensional array object" and utilities and operations built around array objects for fast manipulation. Currently it supports most of the mainstream scientific and machine learning packages available, like `SciPy` and `SciKit Learn`.
- `Kubernetes` [27] is a container orchestration tool initially developed at Google to mimic their internal system Borg. It supports user defined-workloads and maintains a constant system state, ensuring high-availability and reliability.

3 Methodology

In order to effectively and correctly compare the three SBOM standards, we used the tools mentioned in Section 2.2 on each one of the three example projects. This process involved building the tools themselves and running them against each project according to project-specific configuration parameters. The resulting output files were grouped according to each standard, with the filename indicating the tool used for better analysis.

To aid this process, a custom *Python* script was developed. This script is responsible for building each tool, applying the correct flags to the resulting executable and generating an SBOM for each example repository.

This process resulted in 9 reports for CycloneDX, 4 files for SPDX and 0 files for SWID (Even though a sample SWID file was generated outside of this process). A discussion of these results is present in Section 4.

4 Results and Discussion

After running the tool developed by the authors to aid in programmatically generating SBOMs for the example repositories chosen, a total of 13 files have been generated: 9 CycloneDX SBOMs and 4 SPDX SBOMs. No SWID Tag files could be generated using the tool. Nonetheless, a SWID Tag file could be generated outside of the "standard" process. Further discussion is divided per SBOM Standard.

4.1 CycloneDX

When comparing the outputs of the CycloneDX tools (`cdxgen`, `syft`, and `build-info-go`), all produce SBOMs compliant with CycloneDX Version 1.6, enabling easier comparisons. During this analysis, `build-info-go` stopped generating SBOMs for Apache Kafka, so an older SBOM was used.

Among the three tools, `build-info-go` generates the least detailed SBOMs, containing only basic metadata and incomplete component information. For instance, it failed to identify any components or dependencies for *numpy*, likely due to the limited maturity of the BuildInfo format and its conversion to CycloneDX.

`syft` produces the largest SBOMs but lacks completeness in dependency information, as seen with *numpy* and *Kubernetes*. While some dependencies are reported for

Kafka, they appear incomplete. However, **syft** excels in generating extensive component metadata, including licensing, package URLs, CPE entries, and custom properties.

cdxgen, developed by the CycloneDX team, offers superior detail, capturing metadata about the analyzed components and the development lifecycle. It matches **syft** in component metadata while adding features like component hashes, source usages, and sometimes "scope" and evidence fields. Dependency information mirrors **syft**, but **cdxgen** also includes an *annotations* field, suggesting the potential for richer data.

All three tools remain incomplete compared to the comprehensive CycloneDX Object Model, which includes fields like release notes, Service objects, and vulnerability information.

In summary, **build-info-go** is a user-friendly starting point for generating CycloneDX SBOMs. **syft** is a versatile open-source tool with strong SPDX interoperability, while **cdxgen** delivers the most detailed SBOMs. Combining outputs from **syft** and **cdxgen** could enhance SBOM quality. Despite their merits, these tools fall short of the CycloneDX standard's full potential due to the nascent nature of SBOM adoption.

A tool exists for CycloneDX SBOM comparisons but is limited to component-level analysis. Due to repeated patterns in SBOMs, manual analysis proved more effective.

4.2 SPDX

Three tools to generate SPDX SBOMs were tested, giving very different results.

spdx-gradle-plugin was the first tool to be tested, as it is the official SBOM generation tool. As it is Gradle based, we only tested it with the *Kafka* repository. Trying the sample configuration from the project **readme** proved to not be enough, as *Kafka*'s Gradle configuration is really complex, involving several subprojects and intermediary build steps. After some trial and error we were able to configure the plugin for each subproject, generating several SBOMs, but all of them were completely empty. Based on this result we decided to move on to other tools.

spdx-sbom-generator boasted support for all the package managers we tried, but in the end it could only generate an SBOM for *Kubernetes* without issue. Again the problem seems to be the complex build configurations from *Kafka* and *numpy*, which crashed the tool for the former, and could not even detect the package manager for the latter. Overall this tool is still really immature and requires more testing with large projects to be ready for use.

syft ended up being the best tool, as it was the only one to generate all SBOMs successfully.

Comparing the SBOM generated for *Kafka*, the one generated by **syft** is much larger than the one from **spdx-sbom-generator**, seemingly because **syft** is including a reference to each file in the repository, while **spdx-sbom-generator** only includes external dependencies. Anyhow, even when just looking at the dependencies, **syft** includes a much larger amount of metadata.

In conclusion, the SPDX ecosystem leaves a lot to be desired, with the official and dedicated tools we tried having issues. The best tool is unofficial and polyglot, which seems to indicate that other standards are more well supported.

4.3 SWID

Generating SWID tags turned out to be an unfeasible task due to the lack of proper tools and documentation. The tools provided by NIST only support Java projects built with Maven, which is too restrictive since the largest projects use Gradle nowadays, while the tools found on GitHub require implementing a custom portion of code in the project to generate the tags, which is not as straightforward as the other tools and it might be unsuitable for complex projects. In addition to this, the documentation provided by NIST is often not very clear on how to use the tools, which makes it even harder to generate the SWID tags.

Out of all the previously mentioned tools, only `swid-maven-plugin` was able to generate a tag file for [this](#) Maven project. We also tried to use `swid-generator` to generate a SBOM for Kafka, since it is written in Java using Gradle. However, unfortunately this turned out to not be possible for the reasons mentioned above.

5 Conclusion

Out of the 3 analyzed standards, CycloneDX appears to be the most mature and most used one, which can be seen by the existing tool support. Official tools are constantly being developed and improved. Being backed by a large organization (OWASP Foundation) provides CycloneDX with the resources necessary to grow and invest in further developments. On the other end of the spectrum lies SWID, which has the most critical aspects in usage, as testified by our unsuccessful attempts to generate SWID tags and the lack of variety in tools available for this purpose. This is a clear indicator that this standard is not as widely adopted as the others, as also shown in this previous research [11] where only the 10% of the interviewed people actually used SWID. This is a problem that needs to be addressed, as SWID is a useful standard for software identification and its usage should be encouraged by making it easier to generate the tags. Last but not least, SPDX stands somewhere in between the other two standards, being used by one of the biggest open-source projects of today (the Linux Kernel) while also being backed by a famous organization (The Linux Foundation). While it enjoys some level of support from the existing tooling, our approach revealed that there is still a lot to be done by the official bodies developing SPDX and by the users of the standard before it can attain wider adoption in the industry.

Future work could take the methodology and developed tools of this article and apply them to a bigger variety of SBOM generation tools and/or sample repositories (performing more tests would be unfeasible due to the amount of combinations we would have to analyze). Furthermore, and more specifically for the case of SPDX, future work could use the official `cyclonedx-cli` tool to convert CycloneDX SBOMs into SPDX SBOMs and conduct further analysis from there.

References

- [1] Diomidis Spinellis. “Package Management Systems”. In: *IEEE Software* 29.2 (2012). [Accessed 21-10-2024], pp. 84–86. DOI: [10.1109/MS.2012.38](https://doi.org/10.1109/MS.2012.38).
- [2] ISO Central Secretary. *Information technology — IT asset management — Part 2: Software identification tag*. en. Standard ISO/IEC 19770-2:2015. [Accessed 26-11-2024]. Geneva, CH: International Organization for Standardization, 2015. URL: <https://www.iso.org/standard/65666.html>.
- [3] Éamonn Ó Muirí. “Framing software component transparency: Establishing a common software bill of material (SBOM)”. In: *NTIA, Nov 12* (2019). [Accessed 15-10-2024].
- [4] Labs64. *Swid-generator*. <https://github.com/Labs64/swid-generator>. 2020.
- [5] Ravie Lakshmanan. *Extremely Critical Log4J Vulnerability Leaves Much of the Internet at Risk*. [Accessed 27-11-2024]. 2021. URL: <https://thehackernews.com/2021/12/extremely-critical-log4j-vulnerability.html>.
- [6] Sean Peisert et al. “Perspectives on the SolarWinds Incident”. In: *IEEE Security & Privacy* 19.2 (2021), pp. 7–13. DOI: [10.1109/MSEC.2021.3051235](https://doi.org/10.1109/MSEC.2021.3051235).
- [7] Boming Xia et al. “An Empirical Study on Software Bill of Materials: Where We Stand and the Road Ahead”. In: *2023 IEEE/ACM 45th International Conference on Software Engineering (ICSE)*. [Accessed 15-10-2024]. 2023, pp. 2630–2642. DOI: [10.1109/ICSE48619.2023.00219](https://doi.org/10.1109/ICSE48619.2023.00219).
- [8] Nusrat Zahan et al. “Software Bills of Materials Are Required. Are We There Yet?” In: *IEEE Security & Privacy* 21.2 (2023). [Accessed 15-10-2024], pp. 82–88. DOI: [10.1109/MSEC.2023.3237100](https://doi.org/10.1109/MSEC.2023.3237100).
- [9] Tingting Bi et al. “On the Way to SBOMs: Investigating Design Issues and Solutions in Practice”. In: *ACM Trans. Softw. Eng. Methodol.* 33.6 (June 2024). [Accessed 15-10-2024]. ISSN: 1049-331X. DOI: [10.1145/3654442](https://doi.org/10.1145/3654442). URL: <https://doi.org/10.1145/3654442>.
- [10] Berend Kloeg et al. “Charting the Path to SBOM Adoption: A Business Stakeholder-Centric Approach”. In: *Proceedings of the 19th ACM Asia Conference on Computer and Communications Security*. ASIA CCS ’24. [Accessed 15-10-2024]. Singapore, Singapore: Association for Computing Machinery, 2024, pp. 1770–1783. ISBN: 9798400704826. DOI: [10.1145/3634737.3637659](https://doi.org/10.1145/3634737.3637659). URL: <https://doi.org/10.1145/3634737.3637659>.
- [11] Trevor Stalnaker et al. “BOMs Away! Inside the Minds of Stakeholders: A Comprehensive Study of Bills of Materials for Software Systems”. In: *Proceedings of the IEEE/ACM 46th International Conference on Software Engineering*. ICSE ’24. [Accessed 15-10-2024]. Lisbon, Portugal: Association for Computing Machinery, 2024. ISBN: 9798400702174. DOI: [10.1145/3597503.3623347](https://doi.org/10.1145/3597503.3623347). URL: <https://doi.org/10.1145/3597503.3623347>.
- [12] National Institute of Standards and Technology (NIST). *Software Identification (SWID) Tagging — CSRC — CSRC — csrc.nist.gov*. <https://csrc.nist.gov/projects/Software-Identification-SWID>. [Accessed 15-10-2024].

- [13] Open Worldwide Application Security Project (OWASP). *OWASP CycloneDX Software Bill of Materials (SBOM) Standard* — cyclonedx.org. <https://cyclonedx.org/>. [Accessed 15-10-2024].
- [14] CycloneDX Core Team. *cdxgen documentation* — cyclonedx.github.io. [Accessed 24-11-2024]. URL: <https://cyclonedx.github.io/cdxgen/>.
- [15] Apache Foundation. *GitHub - apache/kafka: Mirror of Apache Kafka* — github.com. <https://github.com/apache/kafka>. [Accessed 25-11-2024].
- [16] Ms Reena Garg. *ISO/IEC JTC 1/SC 7*. <https://www.iso.org/committee/45086.html>. [Accessed 26-11-2024].
- [17] *GitHub - anchore/syft: CLI tool and library for generating a Software Bill of Materials from container images and filesystems* — github.com. <https://github.com/anchore/syft>. [Accessed 29-11-2024].
- [18] *GitHub - opensbom-generator/spdx-sbom-generator: Support CI generation of SBOMs via go lang tooling.* — github.com. <https://github.com/opensbom-generator/spdx-sbom-generator>. [Accessed 02-12-2024].
- [19] *GitHub - spdx/spdx-gradle-plugin* — github.com. <https://github.com/spdx/spdx-gradle-plugin>. [Accessed 02-12-2024].
- [20] *GitHub - spdx/spdx-maven-plugin: Plugin for supporting SPDX in a Maven build.* — github.com. <https://github.com/spdx/spdx-maven-plugin>. [Accessed 02-12-2024].
- [21] *JFrog Build Info - Build Info by JFrog* — [buildinfo.org](https://www.buildinfo.org). [Accessed 24-11-2024]. URL: <https://www.buildinfo.org/>.
- [22] Linux Foundation. *SPDX Linux Foundation Projects Site* — spdx.dev. <https://spdx.dev/>. [Accessed 15-10-2024].
- [23] NIST. *swid-maven-plugin*. [Accessed 27-11-2024]. URL: <https://pages.nist.gov/swid-tools/swid-maven-plugin/>.
- [24] NodeJS Team. *NPM: Node Package Manager*. [Accessed 15-10-2024]. URL: <https://www.npmjs.com/>.
- [25] Numpy. *GitHub - numpy/numpy: The fundamental package for scientific computing with Python.* — github.com. <https://github.com/numpy/numpy>. [Accessed 28-11-2024].
- [26] Rust Team. *Cargo*. [Accessed 15-10-2024]. URL: <https://doc.rust-lang.org/cargo/>.
- [27] Kubernetes SIG. *GitHub - kubernetes/kubernetes: Production-Grade Container Scheduling and Management* — github.com. <https://github.com/kubernetes/kubernetes>. [Accessed 28-11-2024].
- [28] SPDX. *SPDX Online Tool* — tools.spdx.org. <https://tools.spdx.org/app/>. [Accessed 02-12-2024].
- [29] Davide Baggio, João Pereira, and Nuno Pereira. “System & Software Security - Software Bills of Materials”. [Accessed 24-11-2024]. N.D.

A Appendix A

The code and data used to perform this work can be found at:

<https://github.com/naapperas/uleiden-sss>

B Appendix B




Figure 1: Diagram of the processes followed in this project.